

CS6120: Lecture 4

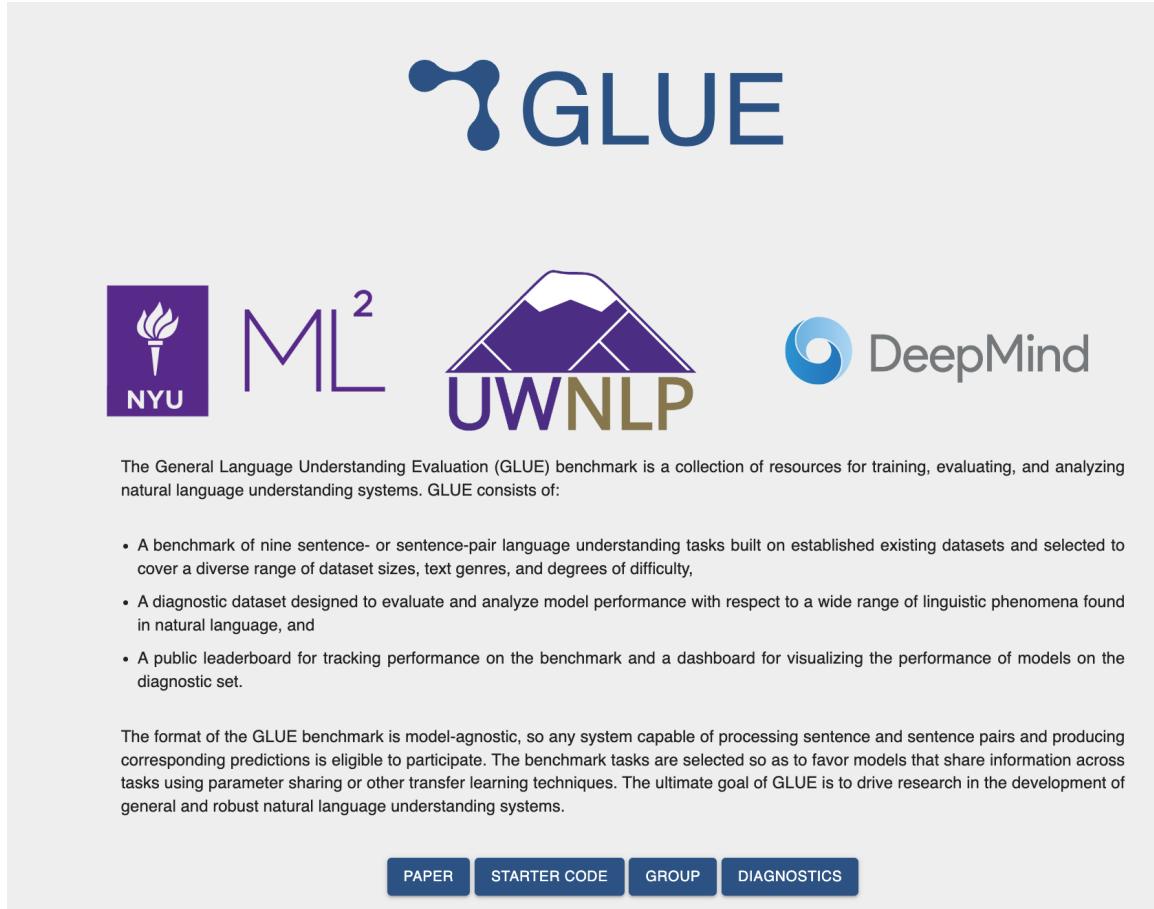
Kenneth Church

<https://kwchurch.github.io>

Agenda

- Homework
 - Assignment 2: [HuggingFace Pipelines](#)
- Background Material
- Old Business
 - [Colab](#)
 - Deep Nets: Inference
 - Classification & Regression
 - Anything → Vector
 - Machine Translation
 - Fill Mask
- New Business
 - [bertviz](#)
 - Deep Nets: Fine-Tuning
 - Easy: inference
 - Hard: pre-training
 - Not too hard: fine-tuning
 - Code: colab
 - [HuggingFace Tutorial](#)
 - [HuggingFace Colab](#)
 - [run glue example](#)

<https://gluebenchmark.com/>



The image shows the GLUE benchmark landing page. At the top center is the GLUE logo, featuring a blue stylized 'G' followed by the word 'GLUE'. Below the logo are logos for NYU (purple square with a torch), ML² (purple 'ML' with a small '2'), UW NLP (purple mountain peak with 'UW NLP' text), and DeepMind (blue circular logo with 'DeepMind'). A descriptive text block follows, detailing the GLUE benchmark's purpose and components. At the bottom are four blue buttons labeled 'PAPER', 'STARTER CODE', 'GROUP', and 'DIAGNOSTICS'.

The General Language Understanding Evaluation (GLUE) benchmark is a collection of resources for training, evaluating, and analyzing natural language understanding systems. GLUE consists of:

- A benchmark of nine sentence- or sentence-pair language understanding tasks built on established existing datasets and selected to cover a diverse range of dataset sizes, text genres, and degrees of difficulty,
- A diagnostic dataset designed to evaluate and analyze model performance with respect to a wide range of linguistic phenomena found in natural language, and
- A public leaderboard for tracking performance on the benchmark and a dashboard for visualizing the performance of models on the diagnostic set.

The format of the GLUE benchmark is model-agnostic, so any system capable of processing sentence and sentence pairs and producing corresponding predictions is eligible to participate. The benchmark tasks are selected so as to favor models that share information across tasks using parameter sharing or other transfer learning techniques. The ultimate goal of GLUE is to drive research in the development of general and robust natural language understanding systems.

PAPER STARTER CODE GROUP DIAGNOSTICS

Some BERT Models

<https://github.com/google-research/bert>

- **BERT-Large, Uncased (Whole Word Masking):**
 - 24-layer, 1024-hidden, 16-heads, 340M parameters
- **BERT-Large, Cased (Whole Word Masking):**
 - 24-layer, 1024-hidden, 16-heads, 340M parameters
- **BERT-Base, Uncased:**
 - 12-layer, 768-hidden, 12-heads, 110M parameters
- **BERT-Large, Uncased:**
 - 24-layer, 1024-hidden, 16-heads, 340M parameters
- **BERT-Base, Cased:**
 - 12-layer, 768-hidden, 12-heads , 110M parameters
- **BERT-Large, Cased:**
 - 24-layer, 1024-hidden, 16-heads, 340M parameters
- **BERT-Base, Multilingual Cased (New, recommended):**
 - 104 languages, 12-layer, 768-hidden, 12-heads, 110M parameters

Smaller Models

<https://github.com/google-research/bert>

	H=128	H=256	H=512	H=768
L=2	2/128 (BERT-Tiny)	2/256	2/512	2/768
L=4	4/128	4/256 (BERT-Mini)	4/512 (BERT-Small)	4/768
L=6	6/128	6/256	6/512	6/768
L=8	8/128	8/256	8/512 (BERT-Medium)	8/768
L=10	10/128	10/256	10/512	10/768
L=12	12/128	12/256	12/512	12/768 (BERT-Base)

Note that the BERT-Base model in this release is included for completeness only; it was re-trained under the same regime as the original model.

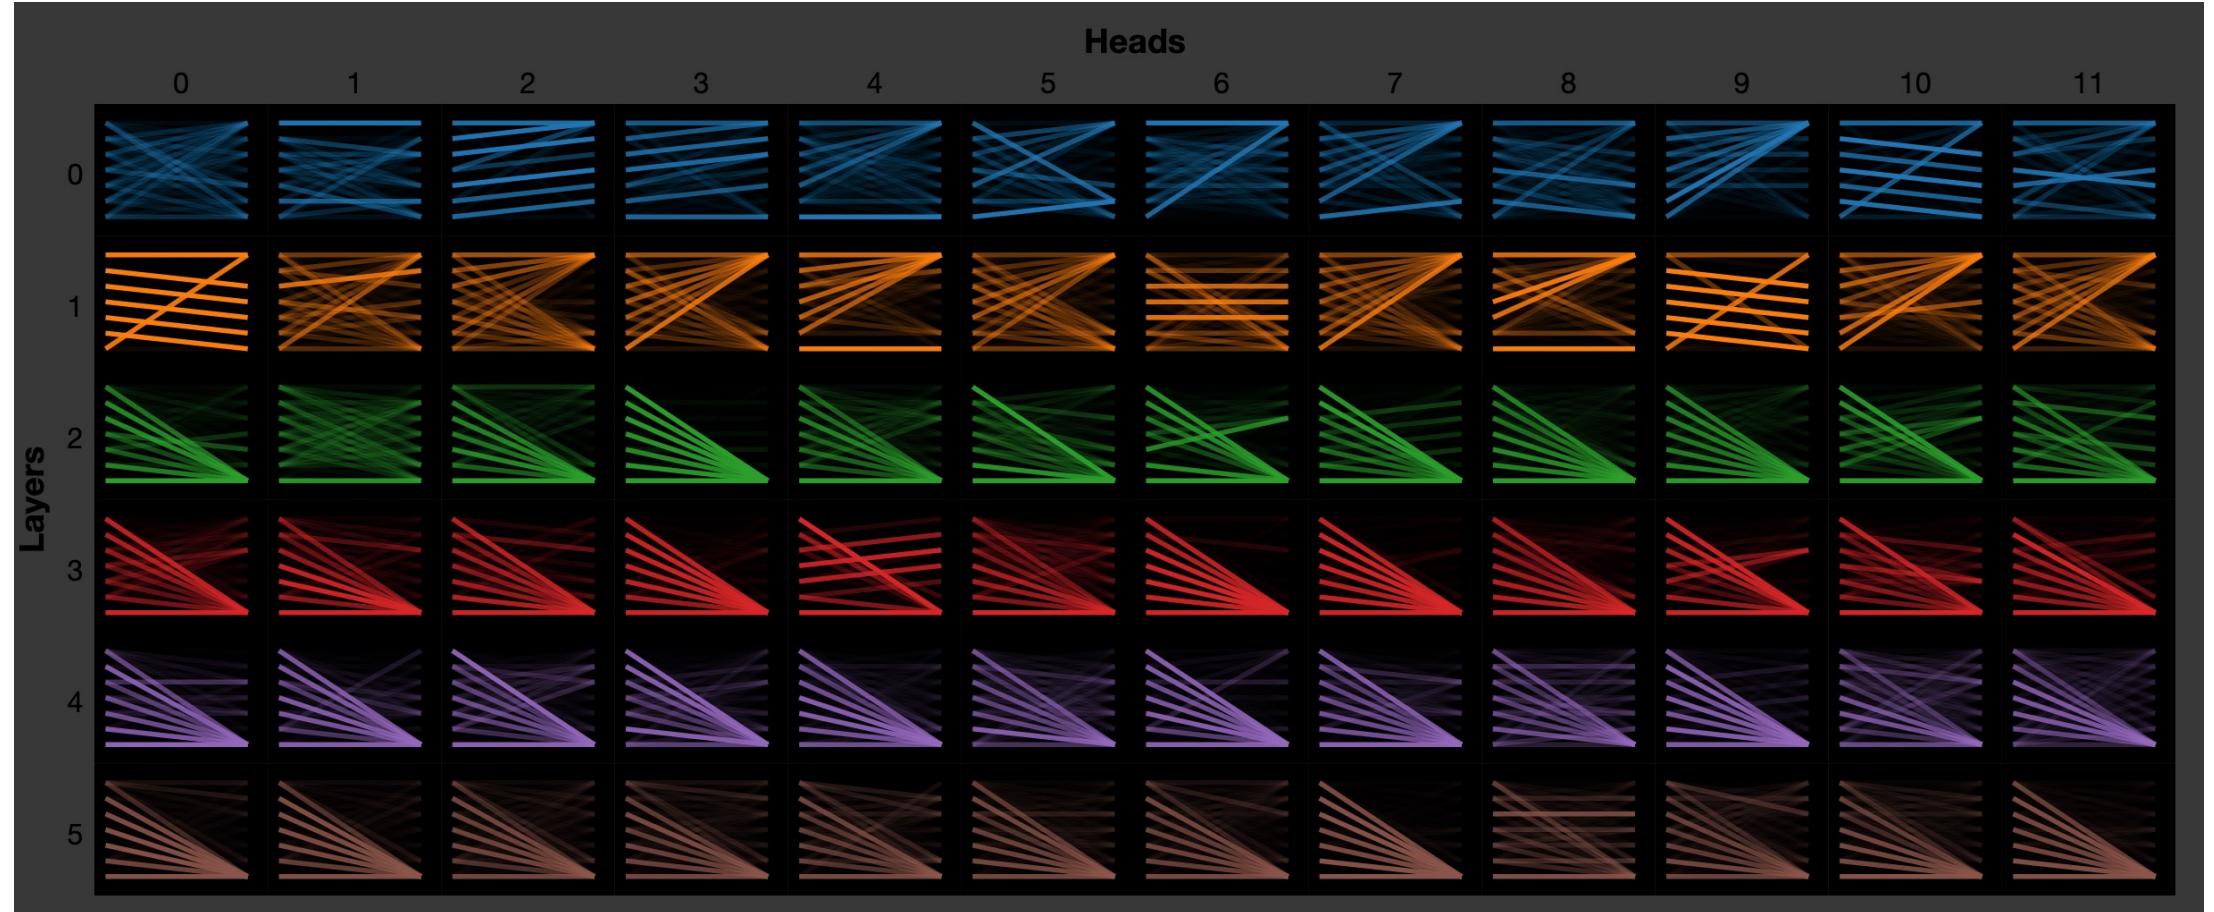
Here are the corresponding GLUE scores on the test set:

Model	Score	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m	MNLI-mm	QNLI(v2)
BERT-Tiny	64.2	0.0	83.2	81.1/71.1	74.3/73.6	62.2/83.4	70.2	70.3	81.5
BERT-Mini	65.8	0.0	85.9	81.1/71.8	75.4/73.3	66.4/86.2	74.8	74.3	84.1
BERT-Small	71.2	27.8	89.7	83.4/76.2	78.8/77.0	68.1/87.0	77.6	77.0	86.4
BERT-Medium	73.5	38.0	89.6	86.6/81.6	80.4/78.4	69.6/87.9	80.0	79.1	87.7

- Trade-off: Costs v. Performance
- Costs:
 - Space
 - Time
 - Power
- Industry Practice
 - Train large models and then make them smaller (distillation)

What do we mean by layers, heads, etc.?

[colab](#)



A Multiscale Visualization of Attention in the Transformer Model (bertviz)

<https://aclanthology.org/P19-3007.pdf>

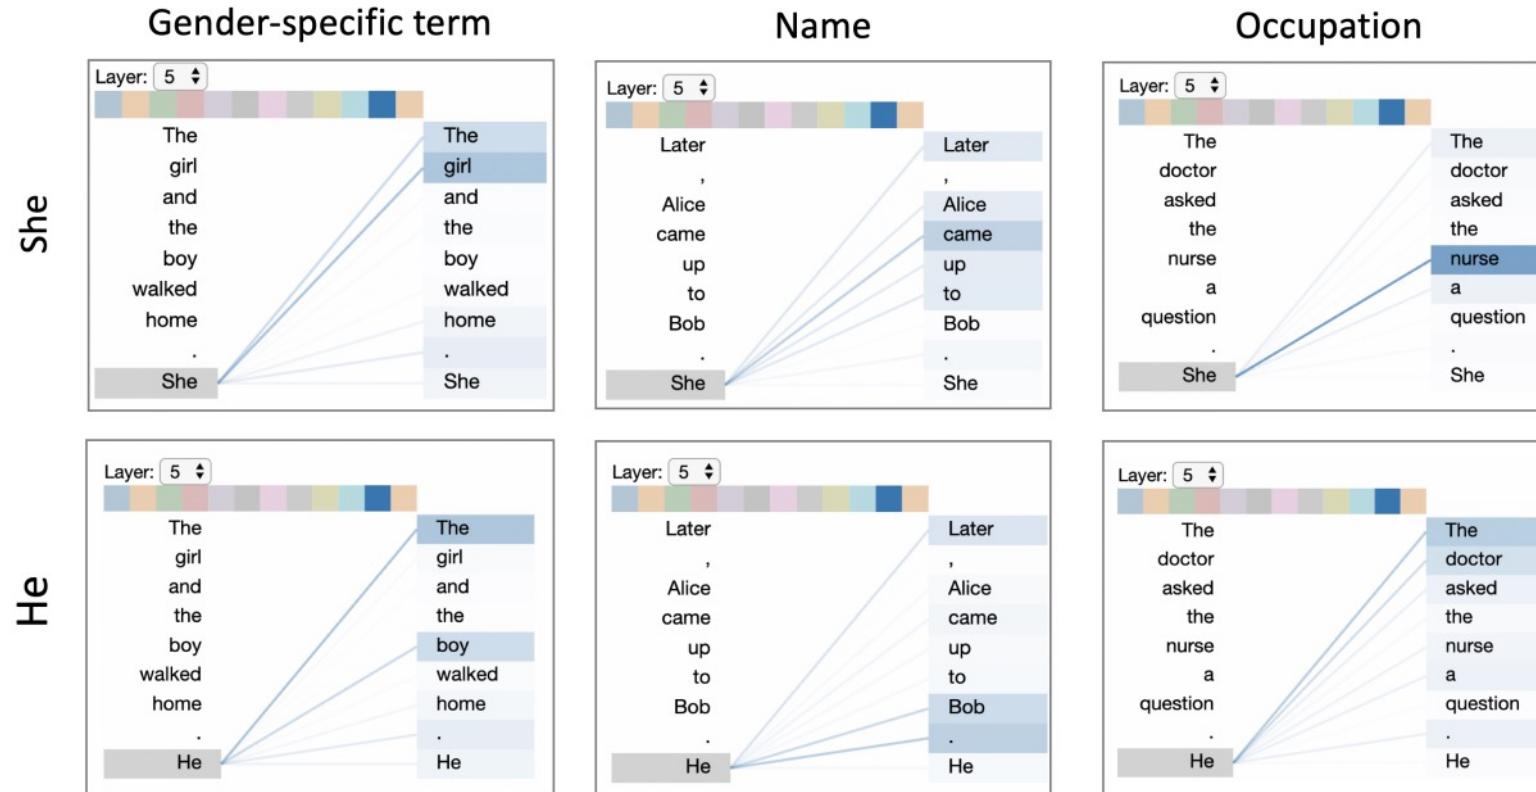


Figure 4: Attention pattern in GPT-2 related to coreference resolution suggests the model may encode gender bias.

<https://arxiv.org/pdf/1906.04341.pdf>

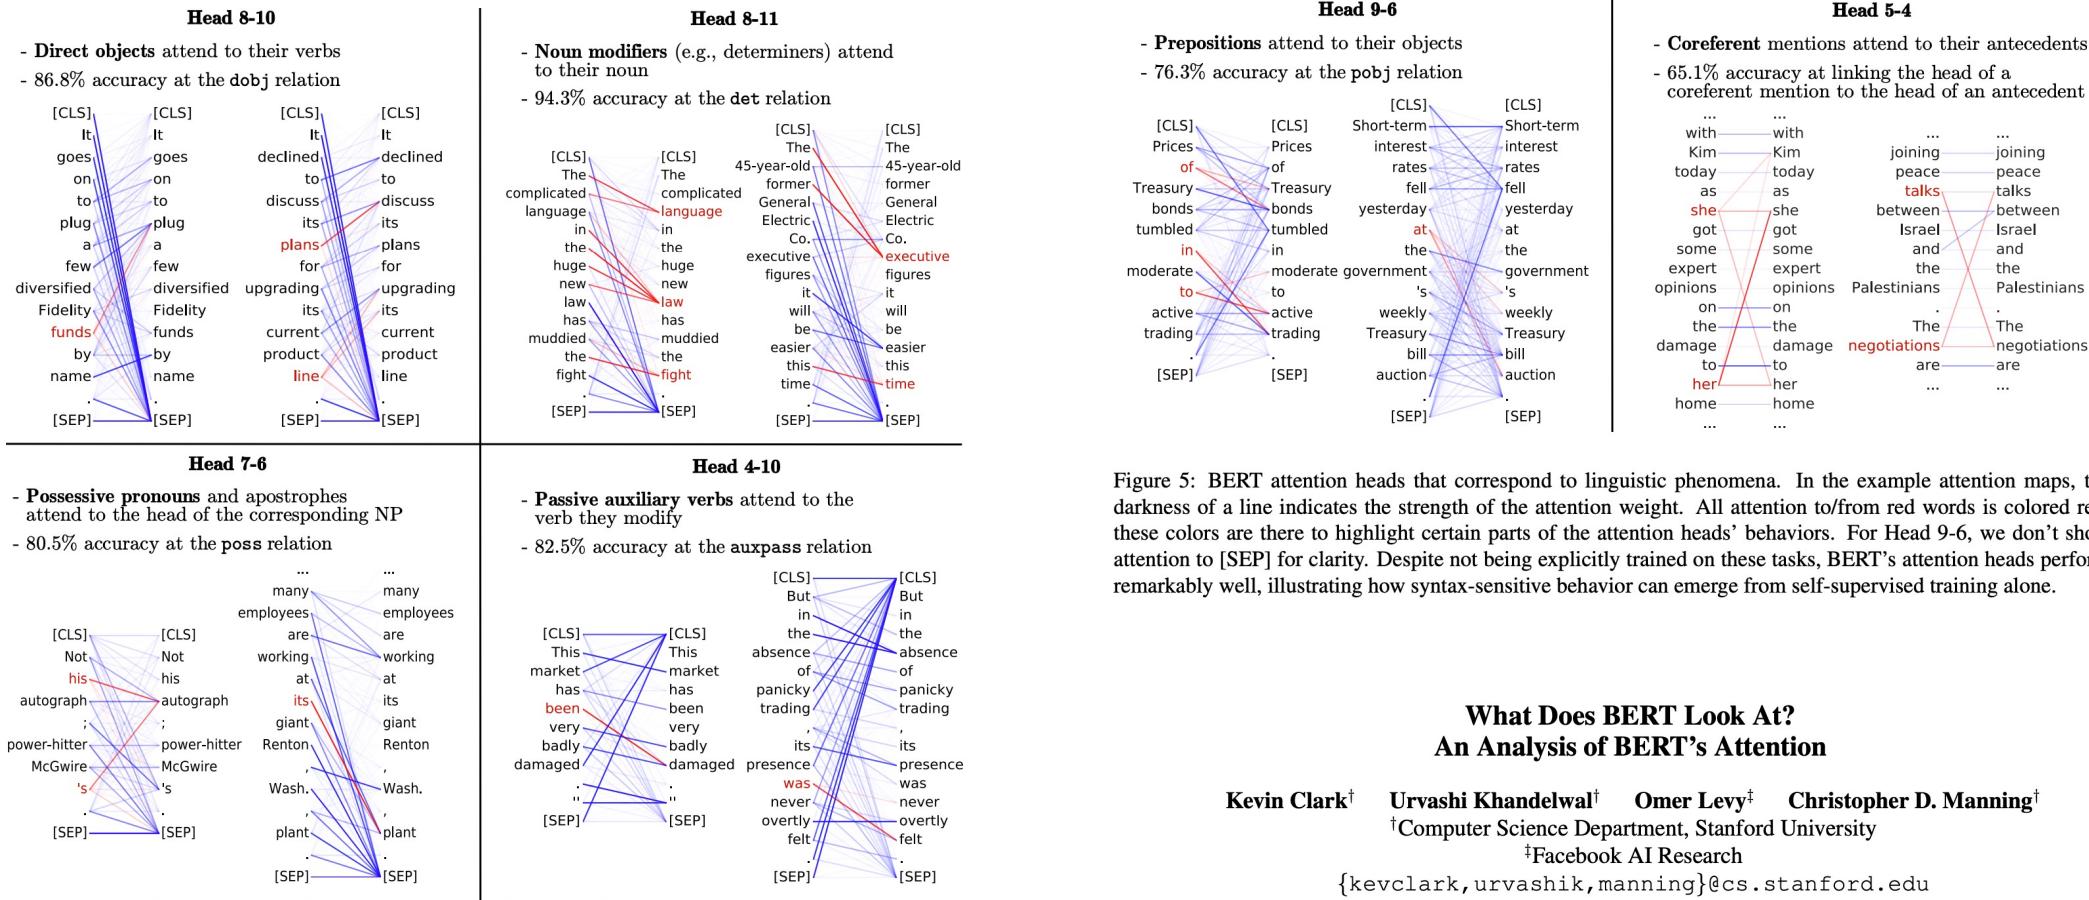


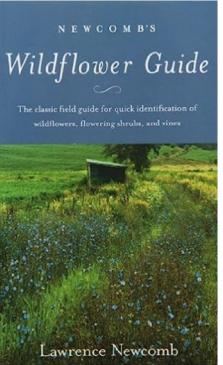
Figure 5: BERT attention heads that correspond to linguistic phenomena. In the example attention maps, the darkness of a line indicates the strength of the attention weight. All attention to/from red words is colored red; these colors are there to highlight certain parts of the attention heads' behaviors. For Head 9-6, we don't show attention to [SEP] for clarity. Despite not being explicitly trained on these tasks, BERT's attention heads perform remarkably well, illustrating how syntax-sensitive behavior can emerge from self-supervised training alone.

Review

Standard 3-Step Recipe

<u>Step</u>	<u>Description</u>	<u>Time</u>	<u>Hardware</u>
1	Pre-Training	Days/Weeks	Large GPU cluster
2	Fine-Tuning (fit)	Hours/Days	1+ GPUs
3	Inference (predict)	Seconds/Minutes	0+ GPUs

Not Hard



Example of Fine-Tuning (aka, *fit*)

***fit*:** $f_{pre} + \text{data} \rightarrow f_{post}$

f_{pre} : Resnet

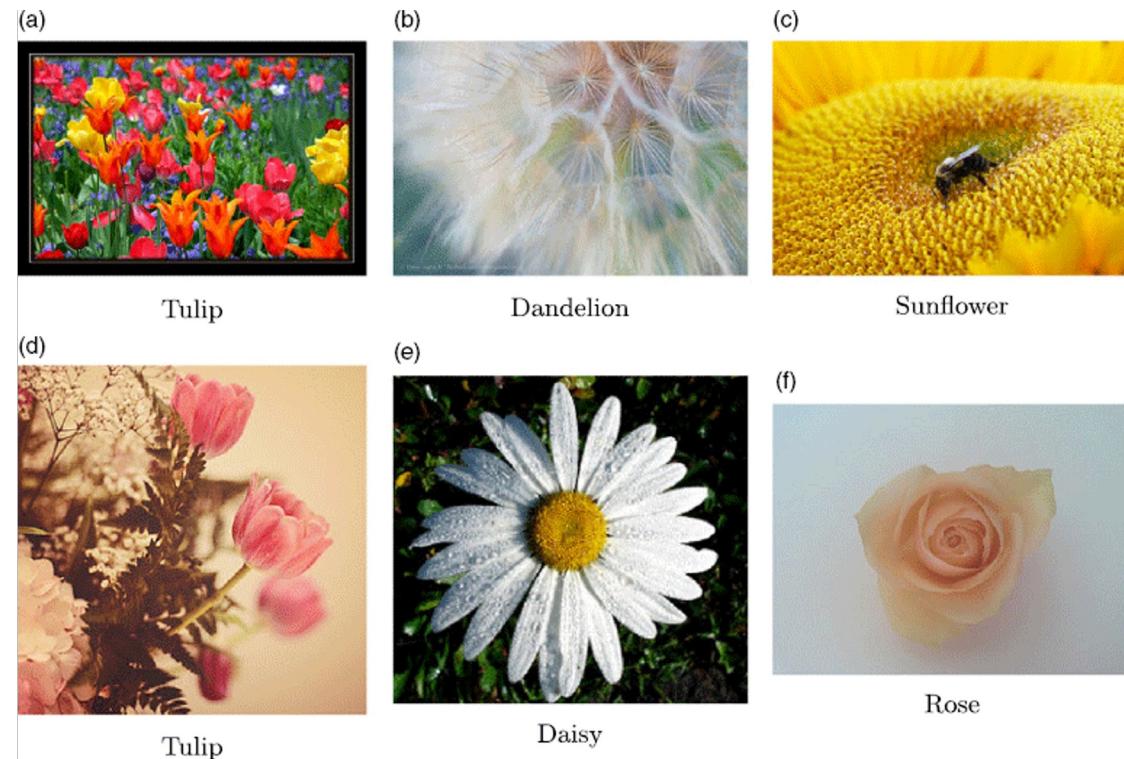
- Maps images (jpg files) → classes (strings)
- Trained on ImageNet
 - input (x): 14M images (of many things)
 - output (y): 1000 classes (strings)

- **data : flowers**

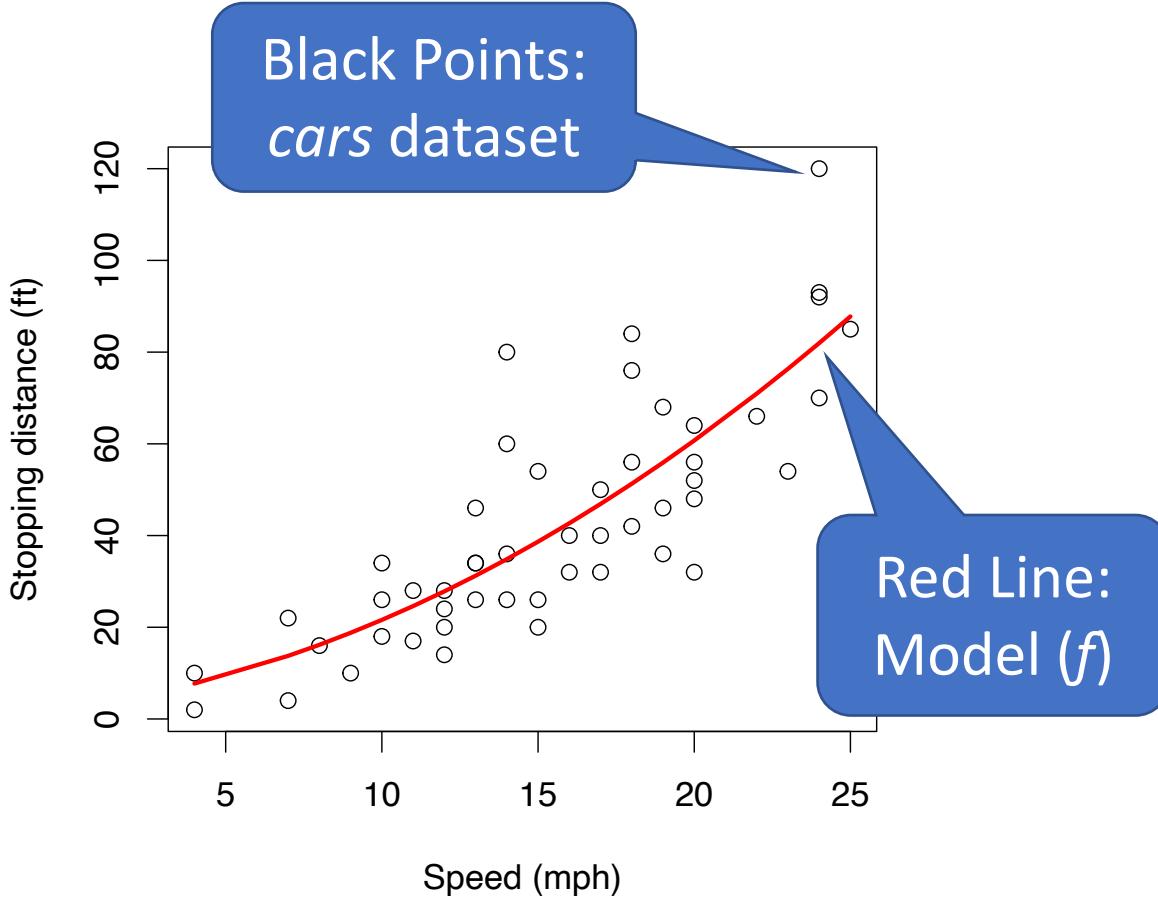
- input (x): 2195 pictures of flowers
- outputs (y): 5 classes of flowers
 - *Tulip, Dandelion, Sunflower, Daisy, Rose*

- f_{post} :

- Maps images (jpg files) → flowers (strings)
- Reject modeling is hard



Fine-Tuning (Fit) in R (Statistics Package)

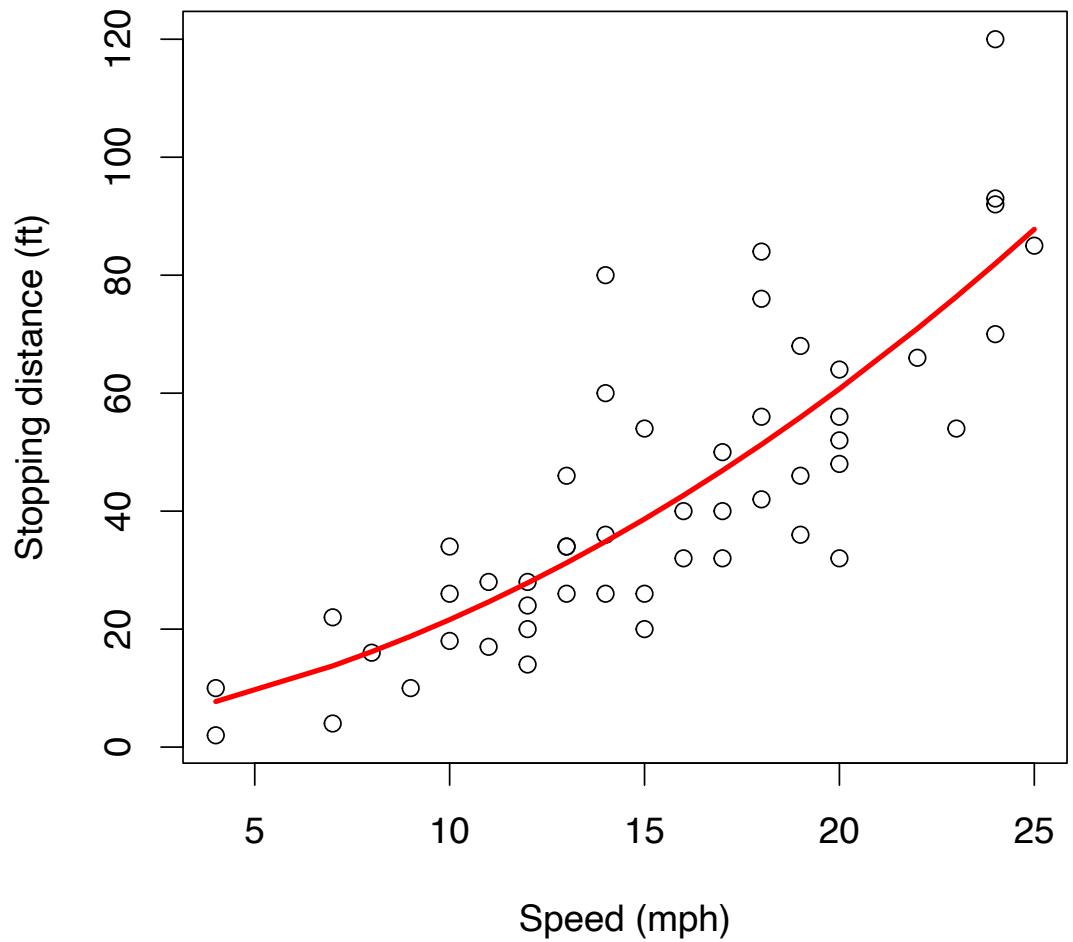


Make deep nets look like regression
(Not much programming)

- Notational Conventions:
 - Observations: Circles
 - Models (f): Red Lines
- Prediction: $f(x)$
 - Use model (f) to map
 - input x (speed) to
 - output y (stopping distance)
 - For linear regression,
 - f is a polynomial
 - For gft,
 - f is typically a model from a hub

Datasets

Example: Cars



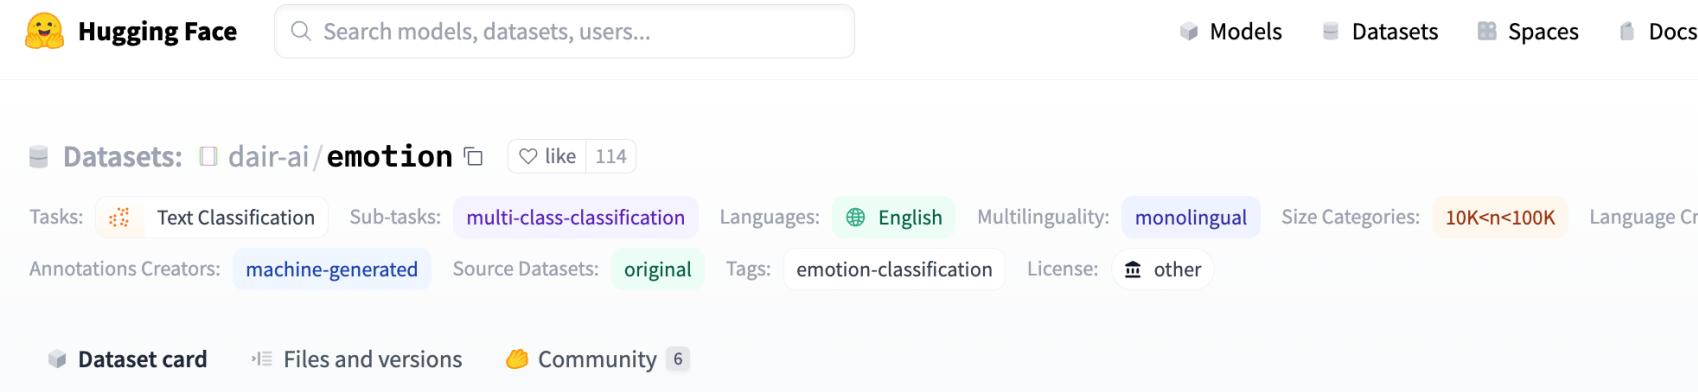
```
> head(cars)
```

	speed	dist
1	4	2
2	4	10
3	7	4
4	7	22
5	8	16
6	9	10

Two Columns:
1. cars\$speed
2. cars\$dist

Example of Datasets in HuggingFace

<https://huggingface.co/datasets/dair-ai/emotion>



Hugging Face Search models, datasets, users... Models Datasets Spaces Docs

Datasets: dair-ai/emotion like 114

Tasks: Text Classification Sub-tasks: multi-class-classification Languages: English Multilinguality: monolingual Size Categories: 10K< n <100K Language Creators:

Annotations Creators: machine-generated Source Datasets: original Tags: emotion-classification License: other

Dataset card Files and versions Community 6

Subset Split

Auto-converted to Parquet API Go to dataset viewer

X Search this dataset

text (string) y label (class label)

"i didnt feel humiliated"	0 (sadness)
"i can go from feeling so hopeless to so damned hopeful just from being around someone who cares and is awake"	0 (sadness)
"im grabbing a minute to post i feel greedy wrong"	3 (anger)
"i am ever feeling nostalgic about the fireplace i will know that it is still on the property"	2 (love)
"i am feeling grouchy"	3 (anger)
"ive been feeling a little burdened lately wasnt sure why that was" 9/11/2023	0 (sadness) CS6120
"ive been taking or milligrams or times recommended amount and ive fallen asleep a lot faster but i also feel like so funny"	5 (surprise)

> head(cars)		
	speed	dist
1	4	2
2	4	10
3	7	4
4	7	22
5	8	16
6	9	10

Downloads last month 7,358

Use in dataset library Edit dataset card
Train in AutoTrain Papers with Code
Evaluate models HF Leaderboard :
...

Homepage: github.com Size of downloaded dataset files: 16.1 MB

Size of the auto-converted Parquet files: 28.2 MB Number of rows: 436,809

Models trained or fine-tuned on dair-ai/emotion

aiknowyou/it-emotion-analyzer
Text Classification Updated Mar 23 1,169 2

Emotion → GLUE (Cola)

<https://huggingface.co/datasets/glue/viewer/cola/train>

The screenshot shows the Hugging Face dataset viewer interface for the 'glue' dataset, specifically the 'cola' subset. The top navigation bar includes links for Models, Datasets, Spaces, Docs, Solutions, and Pricing. Below the navigation, there are filters for Tasks (natural-language-inference, acceptability-classification, text-classification-other-paraphrase-identification), Task Categories (text-classification, text-scoring), Languages (en), and Size Categories (10K < n < 100K). The dataset card shows the 'cola' subset with a preview of 10 rows of data. A blue callout bubble labeled 'Terminology: Subset' points to the 'subset' dropdown menu, which is currently set to 'cola'. Another blue callout bubble labeled 'Terminology: Splits' points to the 'split' dropdown menu, which is currently set to 'train'. The data table has columns for sentence (string), label (class label), and idx (int). The data rows are as follows:

sentence (string)	label (class label)	idx (int)
Our friends won't buy this analysis, let alone the next one we propose.	1 (acceptable)	0
One more pseudo generalization and I'm giving up.	1 (acceptable)	1
One more pseudo generalization or I'm giving up.	1 (acceptable)	2
The more we study verbs, the crazier they get.	1 (acceptable)	3
Day by day the facts are getting murkier.	1 (acceptable)	4
I'll fix you a drink.	1 (acceptable)	5
Fred watered the plants flat.	1 (acceptable)	6
Bill coughed his way out of the restaurant.	1 (acceptable)	7
We're dancing the night away.	1 (acceptable)	8
Herman hammered the metal flat.	1 (acceptable)	9
The critics laughed the play off the stage.	1 (acceptable)	10

>	head(cars)
	speed dist
1	4 2
2	4 10
3	7 4
4	7 22
5	8 16
9	10

Datasets in HuggingFace [colab](#)

CO Datasets in HuggingFace.ipynb ☆

File Edit View Insert Runtime Tools Help Last edited on September 30

Comment Share ⚙️ 

+ Code + Text Connect ▾

Compare with <https://huggingface.co/datasets/imdb> The IMDB dataset has three splits. Each split has two columns (features).

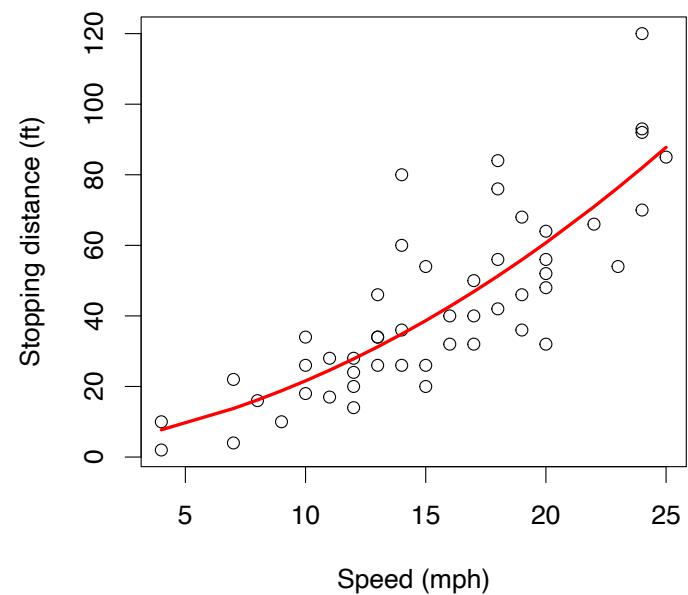
```
{x} [ ] for split in imdb:
    print(split + '\t' + str(imdb[split]))
```

```
train Dataset({
    features: ['text', 'label'],
    num_rows: 25000
})
test Dataset({
    features: ['text', 'label'],
    num_rows: 25000
})
unsupervised Dataset({
    features: ['text', 'label'],
    num_rows: 50000
})
```

```
▶ imdb['train']['text'][0]
```

```
→ I rented I AM CURIOUS-YELLOW from my video store because of all the controversy that surrounded it when it was first released in 1967. I also heard that at first it was seized by U.S. customs if it ever tried to enter this country, therefore being a fan of films considered "controversial" I really had to see this for myself.<br /><br />The plot is centered around a young Swedish drama student named Lena who wants to learn everything she can about life. In particular she wants to focus her attentions to making some sort of documentary on what the average Swede thought about certain political issues such as the Vietnam War and race issues in the United States. In between asking politicians and ordinary denizens of Stockholm about their opinions on politics, she has s
```

glm (General Linear Models) in R (and Sklearn)



```
3 # Create the black points
4 plot(cars, xlab="Speed (mph)", ylab="Stopping distance (ft)")
7 g = glm(dist ~ poly(speed, 2), data=cars)
10 o = order(cars$speed)
11 # Show predictions as a red line
12 lines(cars$speed[o], predict(g,cars)[o], col="red", lwd=3)
```

glm: fit poly model
with data

predict: $dist \approx g(cars\$speed)$

Fine-Tuning (fit) in GFT

$$f_{pre} + data \rightarrow f_{post}$$

- Flowers
 - f_{pre} : Resnet
 - $data$: flowers



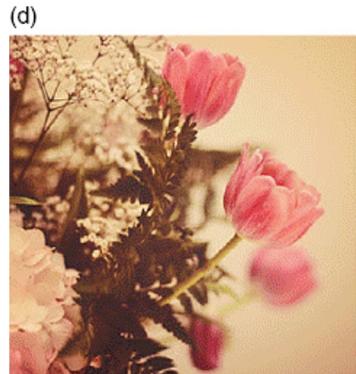
Tulip



Dandelion



Sunflower

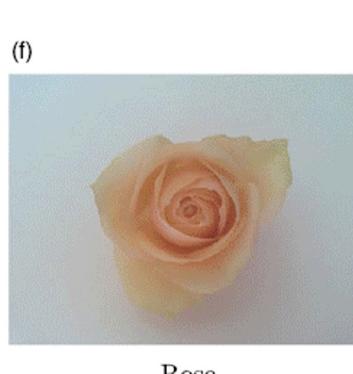


9/11/2023

Tulip



Daisy



Rose

CS6120

```
gft_fit --eqn 'classify: label ~ text' \
--model H:bert-base-cased \
--data H:emotion \
--output_dir $outdir
```

f_{pre} : Pre-trained Model

f_{post} : Post-trained Model

```
7 g = glm(dist ~ poly(speed, 2), data=cars)
```

Fine-Tuning (fit) in R



Datasets: emotion

like 18

Tasks: multi-class-classification text-classification-other-emotion-classification Task Categories: text-classification

Language Creators: machine-generated

Annotations Creators: machine-generated

Source Datasets: original

Dataset card

Files and versions

Dataset Preview

Subset

default

Split

train

text (string)

i didnt feel humiliated

i can go from feeling so hopeless to so damned hopeful just from being around someone who cares and is awake

im grabbing a minute to post i feel greedy wrong

i am ever feeling nostalgic about the fireplace i will know that it is still on the property

i am feeling grouchy

ive been feeling a little burdened lately wasnt sure why that was

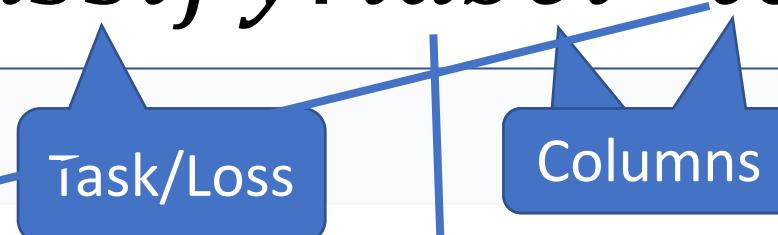
ive been taking or milligrams or times recommended amount and ive fallen asleep a lot faster but i also feel like so funny

i feel as confused about life as a teenager or as jaded as a year old man

i have been with petronas for years i feel that petronas has performed well and made a huge profit

i feel romantic too

i feel like i have to make the suffering i m seeing mean something

Emotion Dataset
classify: label~text

label (class label)

0 (sadness)

0 (sadness)

3 (anger)

2 (love)

3 (anger)

0 (sadness)

5 (surprise)

4 (fear)

1 (joy)

2 (love)

0 (sadness)

Row

Fine-Tuning (fit): Numerous Use Cases

$$f_{pre} + data \rightarrow f_{post}$$

- Flowers
 - f_{pre} : Resnet
 - $data$: flowers
- Emotion Classification
 - f_{pre} : <https://huggingface.co/bert-base-uncase>
 - $data$: <https://huggingface.co/dair-ai/emotion>
- GLUE
- SQuAD
- Machine Translation
- Speech Recognition
- Vision
- and much more

```
gft_fit --eqn 'classify: label ~ text' \
--model H:bert-base-cased \
--data H:emotion \
--output_dir $outdir
```

f_{pre} : Pre-trained Model

f_{post} : Post-trained Model

```
7 g = glm(dist ~ poly(speed, 2), data=cars)
```

Fine-Tuning (fit) in R

GLUE Subsets

Subset	Dataset
COLA	H:glue,cola
SST2	H:glue,sst2
WNLI	H:glue,wnli
MRPC	H:glue,mrpc
QNLI	H:glue,qnli
QQP	H:glue,qqp
SSTB	H:glue,sstb
MNLI	H:glue,mnli

GLUE COLA SUBSET	
Sentence	Label
Bill sang himself to sleep.	1 (acceptable)
Bill squeezed the puppet through the hole.	1 (acceptable)
Bill sang Sue to sleep.	1 (acceptable)
The elevator rumbled itself to the ground.	0 (unacceptable)
If the telephone rang, it could ring itself silly.	1 (acceptable)
She yelled hoarse.	0 (unacceptable)
Ted cried to sleep.	0 (unacceptable)
The tiger bled to death.	1 (acceptable)

GFT Program for COLA

https://github.com/kwchurch/gft/blob/master/examples/fit_examples/model.HuggingFace/language/data.HuggingFace/glue/cola.sh

```
1 #!/bin/sh
2
3 echo hostname = `hostname`
4
5 gft_fit --model H:bert-base-cased \
6   --data H:glue,cola \
7   --metric H:glue,cola \
8   --figure_of_merit matthews_correlation \
9   --output_dir $1 \
10  --eqn 'classify: label ~ sentence' \
11  --num_train_epochs 3
```

GLUE COLA SUBSET

Sentence	Label
Bill sang himself to sleep.	1 (acceptable)
Bill squeezed the puppet through the hole.	1 (acceptable)
Bill sang Sue to sleep.	1 (acceptable)
The elevator rumbled itself to the ground.	0 (unacceptable)
If the telephone rang, it could ring itself silly.	1 (acceptable)
She yelled hoarse.	0 (unacceptable)
Ted cried to sleep.	0 (unacceptable)
The tiger bled to death.	1 (acceptable)

Simple Equations Cover Many Cases of Interest

GLUE: A Popular Benchmark

Subset	Dataset	Equation
COLA	H:glue,cola	<i>classify</i> : $label \sim sentence$
SST2	H:glue,sst2	<i>classify</i> : $label \sim sentence$
WNLI	H:glue,wnli	<i>classify</i> : $label \sim sentence$
MRPC	H:glue,mrpc	<i>classify</i> : $label \sim sentence_1 + sentence_2$
QNLI	H:glue,qnli	<i>classify</i> : $label \sim sentence_1 + sentence_2$
QQP	H:glue,qqp	<i>classify</i> : $label \sim question + sentence$
SSTB	H:glue,sstb	<i>regress</i> : $label \sim question_1 + question_2$
MNLI	H:glue,mnli	<i>classify</i> : $label \sim premise + hypothesis$

Equation Keywords \approx Pipeline Tasks

Benchmark	Subset	Dataset	Equation
GLUE	COLA	H:glue,cola	$classify : label \sim sentence$
SQuAD 1.0		H:squad	$classify_spans: answers \sim question + context$
SQuAD 2.0		H:squad_v2	$classify_spans: answers \sim question + context$
CONLL2003	POS	H:conll2003	$classify_tokens: pos_tags \sim tokens$
	NER	H:conll2003	$classify_tokens: ner_tags \sim tokens$
TIMIT		H:timit_asr	$ctc: text \sim audio$
Amazon Reviews		H:amazon_reviews_multi	$classify : label \sim question + sentence$
VAD		C:\$gft/datasets/VAD/VAD	$regress: Valence + Arousal + Dominance \sim Word$

gft Cheat Sheet (General Fine-Tuning)

4+1 Functions

1. **gft_fit:** $f_{pre} \rightarrow f_{post}$ (fine-tuning)
 - 4 Arguments, --output_dir, --metric, --splits
 - (plus most args in most hubs)
2. **gft_predict:** $f(x) \rightarrow \hat{y}$ (inference)
 - Input: 4 Arguments (x from data or stdin)
 - Output: \hat{y} for each x
3. **gft_eval:** Score model on dataset
 - Input: 4 Arguments, --split, --metric, ...
 - Output: Score
4. **gft_summary:** Find good stuff
 - Input: 4 Arguments
 - (may include: `_contains_`, `_infer_`)
5. **gft_cat_dataset:** Output data to *stdout*
 - Input: 4 Arguments (--data, --eqn)

4 Arguments

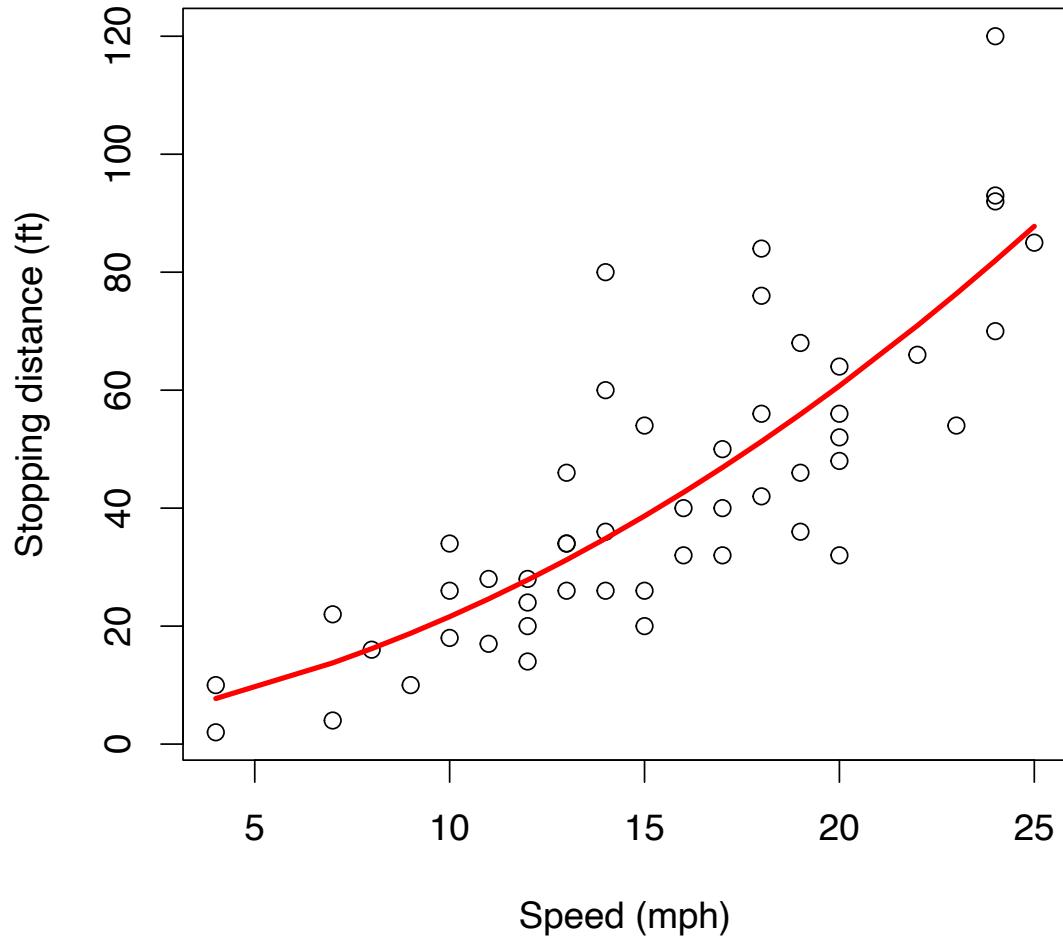
- ✓ --data
- ✓ --model
- --eqn **task:** $y \sim x_1 + x_2$
- --task
 1. classify (text-classification)
 2. classify_tokens (token-classification)
 3. classify_spans (QA, question-answering)
 4. classify_audio (audio-classification)
 5. classify_images (image-classification)
 6. regress
 7. text-generation
 8. MT (translation)
 9. ASR (ctc, automatic-speech-recognition)
 10. fill-mask

Fit a model (g) to data (cars)

Regression in R:

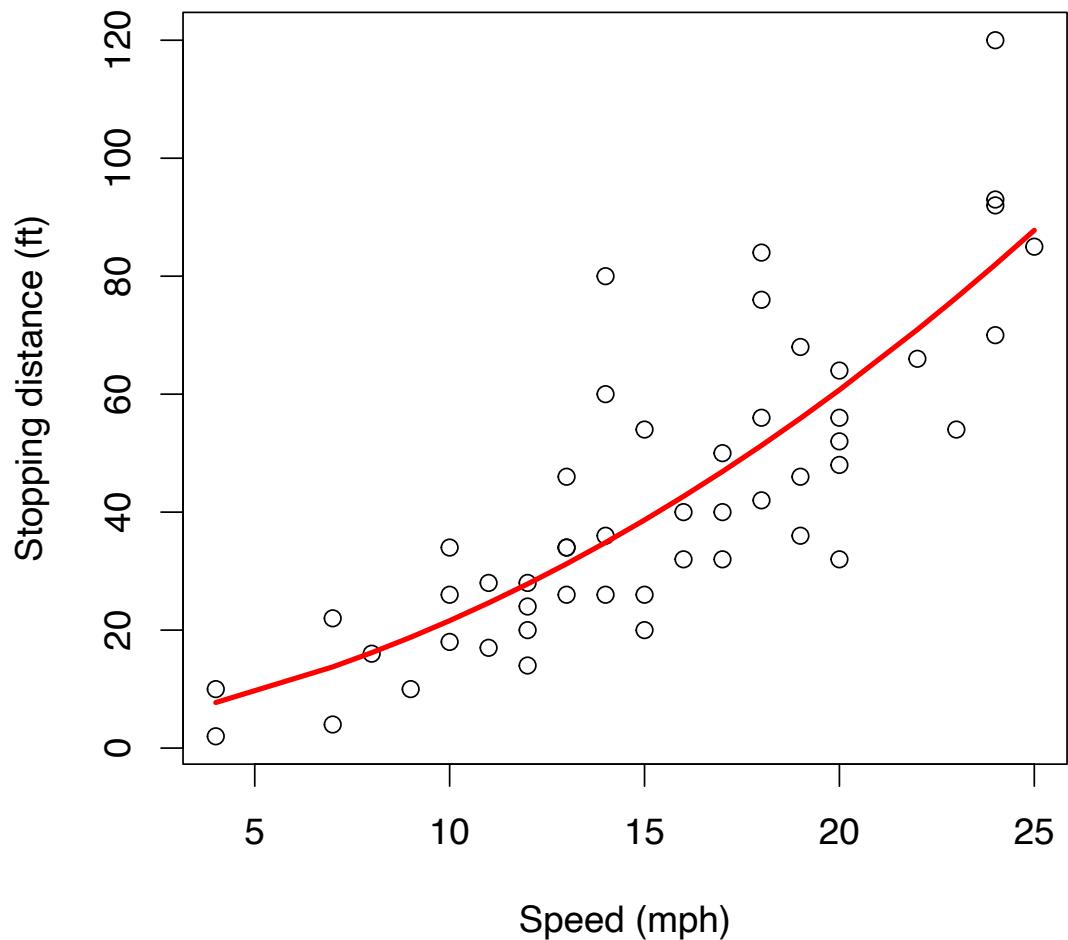
Summarize (almost) anything

```
> plot(cars, xlab="Speed (mph)", ylab="Stopping distance (ft)")  
> g = glm(dist ~ poly(speed,2), data=cars)  
> o = order(cars$speed)  
> lines(cars$speed[o], predict(g,cars)[o], col="red", lwd=3)
```



Regression in R:

Summarize (almost) anything



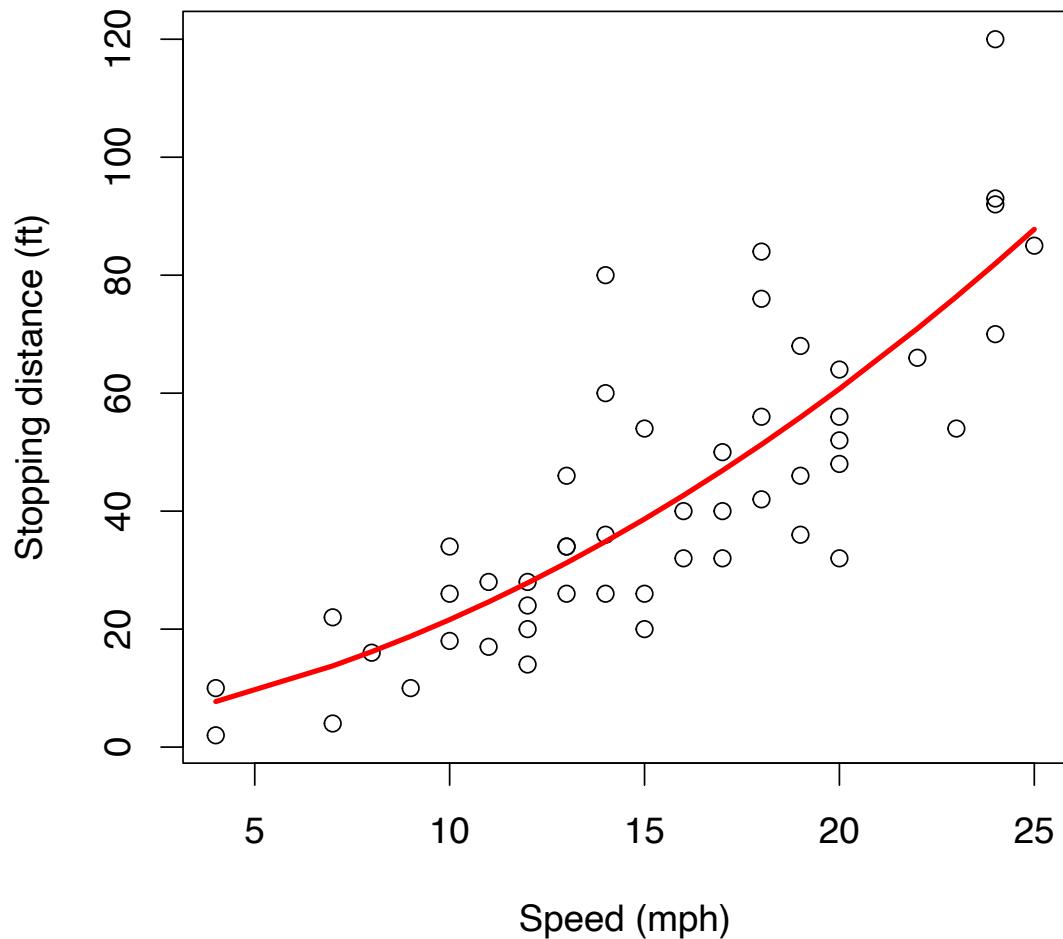
```
> summary(cars)
```

speed	dist
Min. : 4.0	Min. : 2.00
1st Qu.:12.0	1st Qu.: 26.00
Median :15.0	Median : 36.00
Mean :15.4	Mean : 42.98
3rd Qu.:19.0	3rd Qu.: 56.00
Max. :25.0	Max. :120.00

Fit a model (g) to data (cars)

Regression in R:

Summarize (almost) anything



```
> plot(cars, xlab="Speed (mph)", ylab="Stopping distance (ft)")  
> g = glm(dist ~ poly(speed, 2), data=cars)  
> o = order(cars$speed)  
> lines(cars$speed[o], predict(g, cars)[o], col="red", lwd=3)  
> summary(g)
```

Predict

Summarize a model (g)

Call:

```
glm(formula = dist ~ poly(speed, 2), data = cars)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-28.720	-9.184	-3.188	4.628	45.152

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	42.980	2.146	20.026	< 2e-16 ***
poly(speed, 2)1	145.552	15.176	9.591	1.21e-12 ***
poly(speed, 2)2	22.996	15.176	1.515	0.136

Signif. codes: 0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1

(Dispersion parameter for gaussian family taken to be 230.3131)

Null deviance: 32539 on 49 degrees of freedom
Residual deviance: 10825 on 47 degrees of freedom
AIC: 418.77

Number of Fisher Scoring iterations: 2

Opportunity
to Improve g

gft_summary

- Summarize almost anything
 - Models
 - Datasets

Tasks

- classify, text-classification $y \in \{0,1,2, \dots\}$
- regress $y \in \mathbb{R}$ or $y \in \mathbb{R}^N$
- QA, Question Answering, classify spans y for each start/end of span
- token classification
 - NER (Named Entity Recognition)
 - POS (Part of Speech Tagging) y for each token
- translation, MT
- ASR, Automatic Speech Recognition, ctc y for each phoneme

gft Cheat Sheet (General Fine-Tuning)

4+1 Functions

1. **gft_fit:** $f_{pre} \rightarrow f_{post}$ (fine-tuning)
 - 4 Arguments, --output_dir, --metric, --splits
 - (plus most args in most hubs)
2. **gft_predict:** $f(x) \rightarrow \hat{y}$ (inference)
 - Input: 4 Arguments (x from data or stdin)
 - Output: \hat{y} for each x
3. **gft_eval:** Score model on dataset
 - Input: 4 Arguments, --split, --metric, ...
 - Output: Score
4. **gft_summary:** Find good stuff
 - Input: 4 Arguments
 - (may include: `_contains_`, `_infer_`)
5. **gft_cat_dataset:** Output data to *stdout*
 - Input: 4 Arguments (--data, --eqn)

4 Arguments

- ✓ --data
- ✓ --model
- ✓ --eqn *task: $y \sim x_1 + x_2$*
- ✓ --task
 1. classify (text-classification)
 2. classify_tokens (token-classification)
 3. classify_spans (QA, question-answering)
 4. classify_audio (audio-classification)
 5. classify_images (image-classification)
 6. regress
 7. text-generation
 8. MT (translation)
 9. ASR (ctc, automatic-speech-recognition)
 10. fill-mask

HuggingFace run_glue.py colab

GFT

Subset	Dataset	Equation
COLA	H:glue,cola	<i>classify : label ~ sentence</i>
SST2	H:glue,sst2	<i>classify : label ~ sentence</i>
WNLI	H:glue,wnli	<i>classify : label ~ sentence</i>
MRPC	H:glue,mrpc	<i>classify : label ~ sentence1 + sentence2</i>
QNLI	H:glue,qnli	<i>classify : label ~ sentence1 + sentence2</i>
QQP	H:glue,qqp	<i>classify : label ~ question + sentence</i>
SSTB	H:glue,sstb	<i>regress : label ~ question1 + question2</i>
MNLI	H:glue,mnli	<i>classify : label ~ premise + hypothesis</i>

<https://github.com/huggingface/transformers/blob/main/examples/pytorch/text-classification/README.md>

Text classification examples ↗

GLUE tasks ↗

Based on the script `run_glue.py`.

Fine-tuning the library models for sequence classification on the GLUE benchmark: [General Language Understanding Evaluation](#). This script can fine-tune any of the models on the [hub](#) and can also be used for a dataset hosted on our [hub](#) or your own data in a csv or a JSON file (the script might need some tweaks in that case, refer to the comments inside for help).

GLUE is made up of a total of 9 different tasks. Here is how to run the script on one of them:

```
export TASK_NAME=mrpc
python run_glue.py \
--model_name_or_path bert-base-cased \
--task_name $TASK_NAME \
--do_train \
--do_eval \
--max_seq_length 128 \
--per_device_train_batch_size 32 \
--learning_rate 2e-5 \
--num_train_epochs 3 \
--output_dir /tmp/$TASK_NAME/
```

where task name can be one of cola, sst2, mrpc, stsbs, qqp, mnli, qnli, rte, wnli.

HuggingFace run_glue.py [colab](#)

GFT

Subset	Dataset	Equation
COLA	H:glue,cola	<i>classify : label ~ sentence</i>
SST2	H:glue,sst2	<i>classify : label ~ sentence</i>
WNLI	H:glue,wnli	<i>classify : label ~ sentence</i>
MRPC	H:glue,mrpc	<i>classify : label ~ sentence1 + sentence2</i>
QNLI	H:glue,qnli	<i>classify : label ~ sentence1 + sentence2</i>
QQP	H:glue,qqp	<i>classify : label ~ question + sentence</i>
SSTB	H:glue,sstb	<i>regress : label ~ question1 + question2</i>
MNLI	H:glue,mnli	<i>classify : label ~ premise + hypothesis</i>

<https://github.com/huggingface/transformers/blob/main/examples/pytorch/text-classification/README.md>

```
56     task_to_keys = {  
57         "cola": ("sentence", None),  
58         "mnli": ("premise", "hypothesis"),  
59         "mrpc": ("sentence1", "sentence2"),  
60         "qnli": ("question", "sentence"),  
61         "qqp": ("question1", "question2"),  
62         "rte": ("sentence1", "sentence2"),  
63         "sst2": ("sentence", None),  
64         "stsbs": ("sentence1", "sentence2"),  
65         "wnli": ("sentence1", "sentence2"),  
66     }
```

Trainer

<https://huggingface.co/learn/nlp-course/chapter3/3?fw=pt>

```
from datasets import load_dataset
from transformers import AutoTokenizer, DataCollatorWithPadding

raw_datasets = load_dataset("glue", "mrpc")
checkpoint = "bert-base-uncased"
tokenizer = AutoTokenizer.from_pretrained(checkpoint)

def tokenize_function(example):
    return tokenizer(example["sentence1"], example["sentence2"], truncation=True)

tokenized_datasets = raw_datasets.map(tokenize_function, batched=True)
data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

from transformers import AutoModelForSequenceClassification

model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)
```

```
from transformers import TrainingArguments

training_args = TrainingArguments("test-trainer")

from transformers import Trainer

trainer = Trainer(
    model,
    training_args,
    train_dataset=tokenized_datasets["train"],
    eval_dataset=tokenized_datasets["validation"],
    data_collator=data_collator,
    tokenizer=tokenizer,
)

trainer.train()
```

Steps

- Datasets
 - [colab](#)
 - contains splits: train, validation, test
- Tokenizing
 - [colab](#)
 - maps input text to a sequence of tokens
 - where token is an offset into vocabulary
- Model: f
 - Includes tokenizer
 - Typically based on BERT
- Trainer
 - [colab](#)
 - For each epoch (pass over training split)
 - Train (update f , using stochastic gradient descent)
 - Evaluate (score f on validation split)

```
from transformers import TrainingArguments  
  
training_args = TrainingArguments("test-trainer")  
  
from transformers import Trainer  
  
trainer = Trainer(  
    model,  
    training_args,  
    train_dataset=tokenized_datasets["train"],  
    eval_dataset=tokenized_datasets["validation"],  
    data_collator=data_collator,  
    tokenizer=tokenizer,  
)  
  
trainer.train()
```

[https://en.wikipedia.org/wiki/Stochastic%20gradient%20descent](https://en.wikipedia.org/wiki/Stochastic_gradient_descent)

<https://realpython.com/gradient-descent-algorithm-python>

Stochastic gradient descent

Article [Talk](#)

From Wikipedia, the free encyclopedia

Stochastic gradient descent (often abbreviated **SGD**) is an iterative method for optimizing an [objective function](#) with suitable [smoothness](#) properties (e.g. [differentiable](#) or [subdifferentiable](#)). It can be regarded as a [stochastic approximation](#) of [gradient descent](#) optimization, since it replaces the actual gradient (calculated from the entire [data set](#)) by an estimate thereof (calculated from a randomly selected subset of the data). Especially in [high-dimensional](#) optimization problems this reduces the very high [computational burden](#), achieving faster iterations in exchange for a lower convergence rate.^[1]

While the basic idea behind stochastic approximation can be traced back to the [Robbins–Monro algorithm](#) of the 1950s, stochastic gradient descent has become an important optimization method in [machine learning](#).^[2]

Background [\[edit\]](#)

Main article: [M-estimation](#)

See also: [Estimating equation](#)

Both [statistical estimation](#) and [machine learning](#) consider the problem of [minimizing](#) an [objective function](#) that has the form of a sum:

$$Q(w) = \frac{1}{n} \sum_{i=1}^n Q_i(w),$$

where the [parameter](#) w that minimizes $Q(w)$ is to be [estimated](#). Each summand function Q_i is typically associated with the i -th [observation](#) in the [data set](#) (used for training).

https://en.wikipedia.org/wiki/Newton%27s_method

Newton's method

Article Talk

40 languages ▾

Read Edit View history Tools ▾

From Wikipedia, the free encyclopedia

This article is about Newton's method for finding roots. For Newton's method for finding minima, see [Newton's method in optimization](#).

In [numerical analysis](#), **Newton's method**, also known as the **Newton–Raphson method**, named after Isaac Newton and Joseph Raphson, is a root-finding algorithm which produces successively better [approximations](#) to the [roots](#) (or zeroes) of a [real-valued function](#). The most basic version starts with a [real-valued function](#) f , its [derivative](#) f' , and an initial guess x_0 for a [root](#) of f . If f satisfies certain assumptions and the initial guess is close, then

$$x_1 = x_0 - \frac{f(x_0)}{f'(x_0)}$$

is a better approximation of the root than x_0 . Geometrically, $(x_1, 0)$ is the [x-intercept](#) of the [tangent](#) of the [graph](#) of f at $(x_0, f(x_0))$: that is, the improved guess, x_1 , is the unique root of the [linear approximation](#) of f at the initial guess, x_0 . The process is repeated as

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}$$

until a sufficiently precise value is reached. The number of correct digits roughly doubles with each step. This algorithm is first in the class of [Householder's methods](#), succeeded by [Halley's method](#). The method can also be extended to [complex functions](#) and to [systems of equations](#).

Description [edit]

The idea is to start with an initial guess, then to approximate the function by its [tangent line](#), and finally to compute the [x-intercept](#) of this tangent line. This [x-intercept](#) will typically be a better approximation to the original function's root than the first guess, and the method can be [iterated](#).

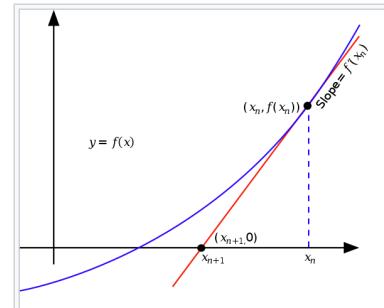
If the [tangent line](#) to the curve $f(x)$ at $x = x_n$ intercepts the [x-axis](#) at x_{n+1} then the slope is

$$f'(x_n) = \frac{f(x_n) - 0}{x_n - x_{n+1}}.$$

Solving for x_{n+1} gives

$$x_{n+1} = x_n - \frac{f(x_n)}{f'(x_n)}.$$

We start the process with some arbitrary initial value x_0 . (The closer to the zero, the better. But, in the absence of any intuition about where the zero might lie, a "guess and check" method might narrow the possibilities to a reasonably small interval by appealing to the [intermediate value theorem](#).) The method will usually converge, provided this initial guess is close enough to the unknown zero, and



https://en.wikipedia.org/wiki/Hill_climbing

Hill climbing

16 languages ▾

Article Talk

Read Edit View history Tools ▾

From Wikipedia, the free encyclopedia

This article is about the mathematical algorithm. For other meanings such as the branch of motorsport, see [Hillclimbing \(disambiguation\)](#).



This article **needs additional citations for verification**. Please help improve this article by adding citations to reliable sources. Unsourced material may be challenged and removed.

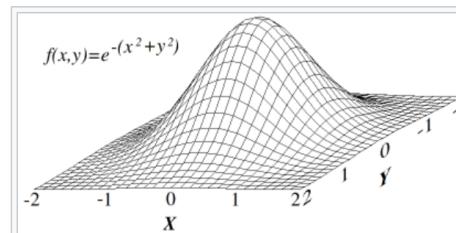
Find sources: "Hill climbing" – news · newspapers · books · scholar · JSTOR (April 2017) ([Learn how and when to remove this template message](#))

In numerical analysis, **hill climbing** is a [mathematical optimization](#) technique which belongs to the family of [local search](#). It is an [iterative algorithm](#) that starts with an arbitrary solution to a problem, then attempts to find a better solution by making an [incremental](#) change to the solution. If the change produces a better solution, another incremental change is made to the new solution, and so on until no further improvements can be found.

For example, hill climbing can be applied to the [travelling salesman problem](#). It is easy to find an initial solution that visits all the cities but will likely be very poor compared to the optimal solution.

The algorithm starts with such a solution and makes small improvements to it, such as switching the order in which two cities are visited. Eventually, a much shorter route is likely to be obtained.

Hill climbing finds optimal solutions for [convex](#) problems – for other problems it will find only [local optima](#) (solutions that cannot be improved upon by any neighboring configurations), which are not necessarily the best possible solution (the [global optimum](#)) out of all possible solutions (the [search space](#)). Examples of algorithms that solve convex problems by hill-climbing include the [simplex algorithm](#) for [linear programming](#) and [binary search](#).^{[1]:253} To attempt to avoid getting stuck in local optima, one could use restarts (i.e. repeated local search), or more complex schemes based on iterations (like [iterated local search](#)), or on memory (like reactive search optimization and [tabu search](#)), or on memory-less stochastic modifications (like [simulated annealing](#)).



A surface with only one maximum. Hill-climbing techniques are well-suited for optimizing over such surfaces, and will converge to the global maximum.

Trainer

<https://huggingface.co/learn/nlp-course/chapter3/3?fw=pt>

```
from datasets import load_dataset
from transformers import AutoTokenizer, DataCollatorWithPadding

raw_datasets = load_dataset("glue", "mrpc")
checkpoint = "bert-base-uncased"
tokenizer = AutoTokenizer.from_pretrained(checkpoint)

def tokenize_function(example):
    return tokenizer(example["sentence1"], example["sentence2"], truncation=True)

tokenized_datasets = raw_datasets.map(tokenize_function, batched=True)
data_collator = DataCollatorWithPadding(tokenizer=tokenizer)

from transformers import AutoModelForSequenceClassification

model = AutoModelForSequenceClassification.from_pretrained(checkpoint, num_labels=2)
```

```
from transformers import TrainingArguments

training_args = TrainingArguments("test-trainer")

from transformers import Trainer

trainer = Trainer(
    model,
    training_args,
    train_dataset=tokenized_datasets["train"],
    eval_dataset=tokenized_datasets["validation"],
    data_collator=data_collator,
    tokenizer=tokenizer,
)

trainer.train()
```

Steps (Review)

- Datasets
 - [colab](#)
 - contains splits: train, validation, test
- Tokenizing
 - [colab](#)
 - maps input text to a sequence of tokens
 - where token is an offset into vocabulary
- Model: f
 - Includes tokenizer
 - Typically based on BERT
- Trainer
 - [colab](#)
 - For each epoch (pass over training split)
 - Train (update f , using stochastic gradient descent)
 - Evaluate (score f on validation split)

```
from transformers import TrainingArguments  
  
training_args = TrainingArguments("test-trainer")  
  
from transformers import Trainer  
  
trainer = Trainer(  
    model,  
    training_args,  
    train_dataset=tokenized_datasets["train"],  
    eval_dataset=tokenized_datasets["validation"],  
    data_collator=data_collator,  
    tokenizer=tokenizer,  
)  
  
trainer.train()
```

Metric v. Loss

- Training Step
 - Optimize loss in SGD
 - Requirement:
 - loss is differentiable
 - Users no longer need to know how to differentiate loss
 - with modern frameworks:
 - pytorch, tensorflow
- Evaluation Step
 - Use metric to score f
 - Ideally, loss \approx metric

Subset	Dataset	Equation
COLA	H:glue,cola	$\text{classify} : \text{label} \sim \text{sentence}$
SST2	H:glue,sst2	$\text{classify} : \text{label} \sim \text{sentence}$
WNLI	H:glue,wnli	$\text{classify} : \text{label} \sim \text{sentence}$
MRPC	H:glue,mrpc	$\text{classify} : \text{label} \sim \text{sentence}_1 + \text{sentence}_2$
QNLI	H:glue,qnli	$\text{classify} : \text{label} \sim \text{sentence}_1 + \text{sentence}_2$
QQP	H:glue,qqp	$\text{classify} : \text{label} \sim \text{question} + \text{sentence}$
SSTB	H:glue,sstb	$\text{regress} : \text{label} \sim \text{question}_1 + \text{question}_2$
MNLI	H:glue,mnli	$\text{classify} : \text{label} \sim \text{premise} + \text{hypothesis}$

Task	Metric	Result	Training time
CoLA	Matthews corr	56.53	3:17
SST-2	Accuracy	92.32	26:06
MRPC	F1/Accuracy	88.85/84.07	2:21
STS-B	Pearson/Spearman corr.	88.64/88.48	2:13
QQP	Accuracy/F1	90.71/87.49	2:22:26
MNLI	Matched acc./Mismatched acc.	83.91/84.10	2:35:23
QNLI	Accuracy	90.66	40:57
RTE	Accuracy	65.70	57
WNLI	Accuracy	56.34	24

<https://github.com/huggingface/transformers/blob/main/examples/pytorch/text-classification/README.md>

Loss Functions in Deep Learning

<https://insideaiml.com/blog/LossFunctions-in-Deep-Learning-1025>

Loss Functions in Deep Learning



Sanober ibrahim
2 years ago

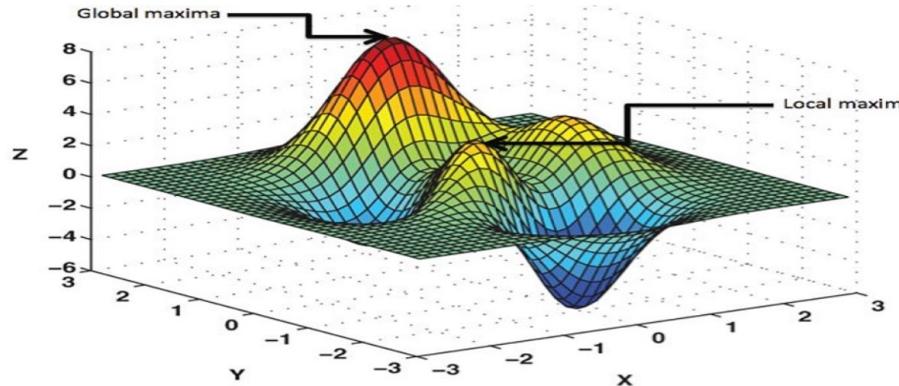


Table of Contents

- What Is a Loss Function?
- Types of Loss Functions
 - 1. Regression Loss Functions
 - 2. Binary Classification Loss Functions
 - 3. Multi-class Classification Loss Functions
- Regression Loss Functions
 - 1. Squared Error Loss
 - L1 and L2 loss
 - 2. Huber Loss
 - 3. Pseudo-Huber loss function
- Binary Classification Loss Functions
 - 1. Hinge Loss
 - 2. Cross-entropy loss
 - 3. Sigmoid-Cross-entropy loss
 - 4. Softmax cross-entropy loss

Winograd Schema (GLUE WNLI)

- The trophy doesn't fit in the brown suitcase
 - because it is too large.
- What is too large?
 - A. The trophy
 - B. The suitcase

Not much better
than chance

Task	Metric	Result	Training time
CoLA	Matthews corr	56.53	3:17
SST-2	Accuracy	92.32	26:06
MRPC	F1/Accuracy	88.85/84.07	2:21
STS-B	Pearson/Spearman corr.	88.64/88.48	2:13
QQP	Accuracy/F1	90.71/87.49	2:22:26
MNLI	Matched acc./Mismatched acc.	83.91/84.10	2:35:23
QNLI	Accuracy	90.66	40:57
RTE	Accuracy	65.70	57
WNLI	Accuracy	56.34	24

Table 1. Time line of the Winograd Schema Challenge.

1972:	Winograd's (1972) thesis introduces the original example.
2010:	Levesque [47] proposes the Winograd Schema Challenge.
2010–2011:	The initial corpus of Winograd schemas is created [50].
2014:	Levesque's Research Excellence talk "On our best behavior" [48].
2016:	The Winograd Schema Challenge is run at IJCAI-16. No systems do much better than chance [16].
2018:	WNLI is incorporated in the GLUE set of benchmarks. BERT-based systems do no better than most-frequent-class guessing [91].
2019, May:	Kocijan et al. [43] achieve 72.5% accuracy on WSC273 using pretraining.
2019, June:	Liu et al. [56] achieve 89.0% on WNLI.
2019, November:	Sakaguchi et al. [77] achieve 90.1% on WSC273.

from: <https://doi.org/10.1016/j.artint.2023.103971>

Winograd Schema (GLUE WNLI)

A Surprisingly Robust Trick for the Winograd Schema Challenge

Vid Kocijan¹, Ana-Maria Crețu², Oana-Maria Camburu^{1,3}, Yordan Yordanov¹, Thomas Lukasiewicz^{1,3}

¹University of Oxford

²Imperial College London

³Alan Turing Institute, London

firstname.lastname@cs.ox.ac.uk, a.cretu@imperial.ac.uk

Abstract

The Winograd Schema Challenge (WSC) dataset Wsc273 and its inference counterpart WNLI are popular benchmarks for natural language understanding and commonsense reasoning. In this paper, we show that the performance of three language models on Wsc273 consistently and robustly improves when fine-tuned on a similar pronoun disambiguation problem dataset (denoted WSCR). We additionally generate a large unsupervised Wsc-like dataset. By fine-tuning the BERT language model both on the introduced and on the WSCR dataset, we achieve overall accuracies of 72.5% and 74.7% on Wsc273 and WNLI, improving the previous state-of-the-art solutions by 8.8% and 9.6%, respectively. Furthermore, our fine-tuned models are also consistently more accurate on the “complex” subsets of Wsc273, introduced by Trichelair et al. (2018).

to the small existing datasets making it difficult to train neural networks directly on the task.

Neural networks have proven highly effective in natural language processing (NLP) tasks, outperforming other machine learning methods and even matching human performance (Hassan et al., 2018; Nangia and Bowman, 2018). However, supervised models require many per-task annotated training examples for a good performance. For tasks with scarce data, transfer learning is often applied (Howard and Ruder, 2018; Johnson and Zhang, 2017), i.e., a model that is already trained on one NLP task is used as a starting point for other NLP tasks.

A common approach to transfer learning in NLP is to train a language model (LM) on large amounts of unsupervised text (Howard and Ruder, 2018) and use it, with or without further fine-tuning, to solve other downstream tasks. Building on top of a LM has proven to be very suc-



Artificial Intelligence

Available online 11 July 2023, 103971

In Press, Corrected Proof

What's this?



The defeat of the Winograd Schema Challenge

Vid Kocijan^{a,1}, Ernest Davis^b, Thomas Lukasiewicz^{c,d}, Gary Marcus^e, Leora Morgenstern^f

Show more ▾

+ Add to Mendeley Share Cite

<https://doi.org/10.1016/j.artint.2023.103971>

Get rights and content

Abstract

The Winograd Schema Challenge—a set of twin sentences involving pronoun reference disambiguation that seem to require the use of commonsense knowledge—was proposed by Hector Levesque in 2011. By 2019, a number of AI systems, based on large pre-trained transformer-based language models and fine-tuned on these kinds of problems, achieved better than 90% accuracy. In this paper, we review the history of the Winograd Schema Challenge and discuss the lasting contributions of the flurry of research that has taken place on the WSC in the last decade. We discuss the significance of various datasets developed for WSC, and the research community’s deeper understanding of the role of surrogate tasks in assessing the intelligence of an AI system.

Keywords

Commonsense reasoning; Winograd Schema Challenge

Steps (Review, again)

- Datasets
 - [colab](#)
 - contains splits: train, validation, test
- Tokenizing
 - [colab](#)
 - maps input text to a sequence of tokens
 - where token is an offset into vocabulary
- Model: f
 - Includes tokenizer
 - Typically based on BERT
- Trainer
 - [colab](#)
 - For each epoch (pass over training split)
 - Train (update f , using stochastic gradient descent)
 - Evaluate (score f on validation split)

```
from transformers import TrainingArguments  
  
training_args = TrainingArguments("test-trainer")  
  
from transformers import Trainer  
  
trainer = Trainer(  
    model,  
    training_args,  
    train_dataset=tokenized_datasets["train"],  
    eval_dataset=tokenized_datasets["validation"],  
    data_collator=data_collator,  
    tokenizer=tokenizer,  
)  
  
trainer.train()
```

https://github.com/huggingface/transformers/blob/main/examples/pytorch/text-classification/run_glue.py

- Trainer

- [colab](#)
- For each epoch
 - Train (update f , using SGD)
 - Evaluate (score f on validation split)

```
550 # Training
551 if training_args.do_train:
552     checkpoint = None
553     if training_args.resume_from_checkpoint is not None:
554         checkpoint = training_args.resume_from_checkpoint
555     elif last_checkpoint is not None:
556         checkpoint = last_checkpoint
557     train_result = trainer.train(resume_from_checkpoint=checkpoint)
558     metrics = train_result.metrics
559     max_train_samples = (
560         data_args.max_train_samples if data_args.max_train_samples is not None else len(train_dataset)
561     )
562     metrics["train_samples"] = min(max_train_samples, len(train_dataset))
563
564     trainer.save_model() # Saves the tokenizer too for easy upload
565
566     trainer.log_metrics("train", metrics)
567     trainer.save_metrics("train", metrics)
568     trainer.save_state()
```

```
570     # Evaluation
571     if training_args.do_eval:
572         logger.info("### Evaluate ***")
573
574     # Loop to handle MNLI double evaluation (matched, mis-matched)
575     tasks = [data_args.task_name]
576     eval_datasets = [eval_dataset]
577     if data_args.task_name == "mnli":
578         tasks.append("mnli-mm")
579         valid_mm_dataset = raw_datasets["validation_mismatched"]
580         if data_args.max_eval_samples is not None:
581             max_eval_samples = min(len(valid_mm_dataset), data_args.max_eval_samples)
582             valid_mm_dataset = valid_mm_dataset.select(range(max_eval_samples))
583         eval_datasets.append(valid_mm_dataset)
584     combined = {}
585
586     for eval_dataset, task in zip(eval_datasets, tasks):
587         metrics = trainer.evaluate(eval_dataset=eval_dataset)
588
589         max_eval_samples = (
590             data_args.max_eval_samples if data_args.max_eval_samples is not None else len(eval_dataset)
591         )
592         metrics["eval_samples"] = min(max_eval_samples, len(eval_dataset))
593
594         if task == "mnli-mm":
595             metrics = {k + "_mm": v for k, v in metrics.items()}
596         if task is not None and "mnli" in task:
597             combined.update(metrics)
598
599     trainer.log_metrics("eval", metrics)
600     trainer.save_metrics("eval", combined if task is not None and "mnli" in task else metrics)
```

https://github.com/huggingface/transformers/blob/main/examples/pytorch/text-classification/run_glue_no_trainer.py

- with trainer
 - 649 lines of code
- without trainer
 - 665 lines of code
 - SGD is more obvious
- in both cases,
 - the example is hardwired for GLUE
- no clear distinction between
 - foreground:
 - GLUE-specific functionality
 - Names of tasks, loss, metric
 - background:
 - General purpose functionality that applies to many/most classification problems

```
539     for epoch in range(starting_epoch, args.num_train_epochs):
540         model.train()
541         if args.with_tracking:
542             total_loss = 0
543             if args.resume_from_checkpoint and epoch == starting_epoch and resume_step is not None:
544                 # We skip the first `n` batches in the dataloader when resuming from a checkpoint
545                 active_dataloader = accelerator.skip_first_batches(train_dataloader, resume_step)
546             else:
547                 active_dataloader = train_dataloader
548             for step, batch in enumerate(active_dataloader):
549                 outputs = model(**batch)
550                 loss = outputs.loss
551                 # We keep track of the loss at each epoch
552                 if args.with_tracking:
553                     total_loss += loss.detach().float()
554                 loss = loss / args.gradient_accumulation_steps
555                 accelerator.backward(loss)
556                 if step % args.gradient_accumulation_steps == 0 or step == len(train_dataloader) - 1:
557                     optimizer.step()
558                     lr_scheduler.step()
559                     optimizer.zero_grad()
560                     progress_bar.update(1)
561                     completed_steps += 1
562
563                     if isinstance(checkpointing_steps, int):
564                         if completed_steps % checkpointing_steps == 0:
565                             output_dir = f"step_{completed_steps}"
566                             if args.output_dir is not None:
567                                 output_dir = os.path.join(args.output_dir, output_dir)
568                             accelerator.save_state(output_dir)
569
570                     if completed_steps >= args.max_train_steps:
571                         break
572
```

https://github.com/huggingface/transformers/blob/main/examples/pytorch/text-classification/run_glue_no_trainer.py

```
539     for epoch in range(starting_epoch, args.num_train_epochs):
540         model.train()
541         if args.with_tracking:
542             total_loss = 0
543         if args.resume_from_checkpoint and epoch == starting_epoch and resume_step is not None:
544             # We skip the first `n` batches in the dataloader when resuming from a checkpoint
545             active_dataloader = accelerator.skip_first_batches(train_dataloader, resume_step)
546         else:
547             active_dataloader = train_dataloader
548         for step, batch in enumerate(active_dataloader):
549             outputs = model(**batch)
550             loss = outputs.loss
551             # We keep track of the loss at each epoch
552             if args.with_tracking:
553                 total_loss += loss.detach().float()
554             loss = loss / args.gradient_accumulation_steps
555             accelerator.backward(loss)
556             if step % args.gradient_accumulation_steps == 0 or step == len(train_dataloader) - 1:
557                 optimizer.step()
558                 lr_scheduler.step()
559                 optimizer.zero_grad()
560                 progress_bar.update(1)
561                 completed_steps += 1
562
563                 if isinstance(checkpointing_steps, int):
564                     if completed_steps % checkpointing_steps == 0:
565                         output_dir = f"step_{completed_steps}"
566                         if args.output_dir is not None:
567                             output_dir = os.path.join(args.output_dir, output_dir)
568                         accelerator.save_state(output_dir)
569
570                 if completed_steps >= args.max_train_steps:
571                     break
572
573         model.eval()
574         samples_seen = 0
575         for step, batch in enumerate(eval_dataloader):
576             with torch.no_grad():
577                 outputs = model(**batch)
578                 predictions = outputs.logits.argmax(dim=-1) if not is_regression else outputs.logits.squeeze()
579                 predictions, references = accelerator.gather((predictions, batch["labels"]))
580                 # If we are in a multiprocess environment, the last batch has duplicates
581                 if accelerator.num_processes > 1:
582                     if step == len(eval_dataloader) - 1:
583                         predictions = predictions[: len(eval_dataloader.dataset) - samples_seen]
584                         references = references[: len(eval_dataloader.dataset) - samples_seen]
585                     else:
586                         samples_seen += references.shape[0]
587                     metric.add_batch(
588                         predictions=predictions,
589                         references=references,
590                     )
591
592         eval_metric = metric.compute()
593         logger.info(f"epoch {epoch}: {eval_metric}")
594
595         if args.with_tracking:
596             accelerator.log(
597                 {
598                     "accuracy" if args.task_name is not None else "glue": eval_metric,
599                     "train_loss": total_loss.item() / len(train_dataloader),
600                     "epoch": epoch,
601                     "step": completed_steps,
602                 },
603                 step=completed_steps,
604             )
605
```

Running Cola on Discovery

Virtual Environment

```
module load python/3.8.1
python3 -m pip install --user --upgrade pip
python3 -m pip install --user virtualenv
python3 -m venv $HOME/venv/transformers
source $HOME/venv/transformers/bin/activate
pip install --upgrade pip
cd /work/k.church/githubs/transformers
pip install -e .

cd /work/k.church/githubs/transformers/examples/pytorch/text-classification/
pip install -r requirements.txt
pip install urllib3==1.26.6
```

Running Cola on Discovery

```
outdir=/courses/CS6120.202410/data/GLUE/model_outputs/cola  
mkdir -p $outdir  
dir=/work/k.church/githubs/transformers/examples/pytorch/text-classification  
cd $outdir  
sh ./cola.sh
```

<https://github.com/google-research/bert>

Model	Score	CoLA	SST-2	MRPC	STS-B	QQP	MNLI-m	MNLI-mm	QNLI(v2)
BERT-Tiny	64.2	0.0	83.2	81.1/71.1	74.3/73.6	62.2/83.4	70.2	70.3	81.5
BERT-Mini	65.8	0.0	85.9	81.1/71.8	75.4/73.3	66.4/86.2	74.8	74.3	84.1
BERT-Small	71.2	27.8	89.7	83.4/76.2	78.8/77.0	68.1/87.0	77.6	77.0	86.4
BERT-Medium	73.5	38.0	89.6	86.6/81.6	80.4/78.4	69.6/87.9	80.0	79.1	87.7

```
d_model/specified_tokens_map.json  
***** train metrics *****  
epoch = 1.0  
train_loss = 0.5148  
train_runtime = 2:18:37.50  
train_samples = 8551  
train_samples_per_second = 1.028  
train_steps_per_second = 0.032  
10/02/2023 15:20:02 - INFO - __main__ - *** Evaluate ***  
[INFO]trainer.py:761] 2023-10-02 15:20:02,237 >> The following columns in the evaluation set don't have a corresponding argument in `BertForSequenceClassification.forward` and have been ignored: sentence, idx. If sentence, idx are not expected by `BertForSequenceClassification.forward`, you can safely ignore this message.  
[INFO]trainer.py:3213] 2023-10-02 15:20:02,240 >> ***** Running Evaluation *****  
[INFO]trainer.py:3215] 2023-10-02 15:20:02,240 >> Num examples = 1043  
[INFO]trainer.py:3218] 2023-10-02 15:20:02,241 >> Batch size = 8  
100% 131/131 [05:03<00:00, 2.32s/it]  
***** eval metrics *****  
epoch = 1.0  
eval_loss = 0.4678  
eval_matthews_correlation = 0.4939  
eval_runtime = 0:05:05.78  
eval_samples = 1043  
eval_samples_per_second = 3.411  
eval_steps_per_second = 0.428  
(transformers) [k.church@c0170 text-classification]$ ls -lt | head  
total 120  
drwxrwx---+ 2 k.church users 4096 Oct 2 15:25 my_colab_model  
-rwxrwx---+ 1 k.church users 18621 Oct 2 12:01 run_xnli.py  
-rwxrwx---+ 1 k.church users 29229 Oct 2 12:01 run_glue_no_trainer.py  
-rwxrwx---+ 1 k.church users 28270 Oct 2 12:01 run_glue.py  
-rwxrwx---+ 1 k.church users 33590 Oct 2 12:01 run_classification.py  
-rw-rw----+ 1 k.church users 112 Oct 2 12:01 requirements.txt  
-rw-rw----+ 1 k.church users 10725 Oct 2 12:01 README.md  
(transformers) [k.church@c0170 text-classification]$ find my_colab_model  
my_colab_model  
my_colab_model/README.md  
my_colab_model/all_results.json  
my_colab_model/config.json  
my_colab_model/eval_results.json  
my_colab_model/pytorch_model.bin
```

Agenda

- Homework
 - Assignment 2: [HuggingFace Pipelines](#)
- Background Material
- Old Business
 - [Colab](#)
 - Deep Nets: Inference
 - Classification & Regression
 - Anything → Vector
 - Machine Translation
 - Fill Mask
- New Business
 - [bertviz](#)
 - Deep Nets: Fine-Tuning
 - Easy: inference
 - Hard: pre-training
 - Not too hard: fine-tuning
 - Code: colab
 - [HuggingFace Tutorial](#)
 - [HuggingFace Colab](#)
 - [run glue example](#)